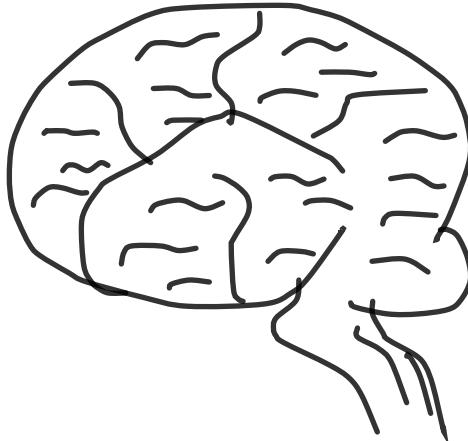
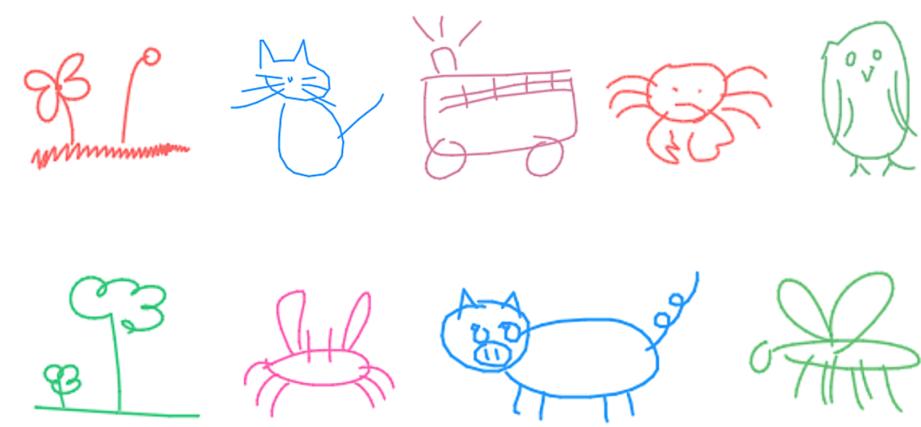
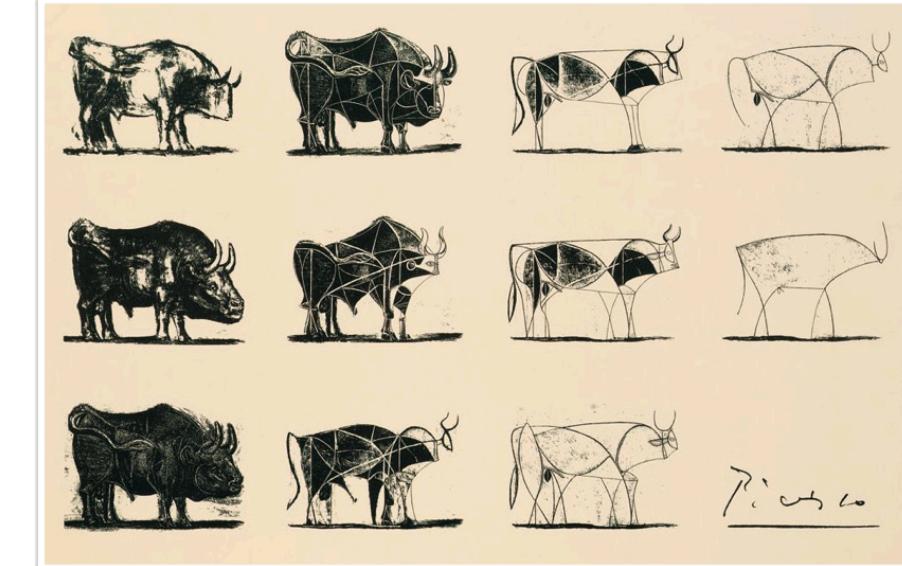


A Neural Representation of Sketch Drawings

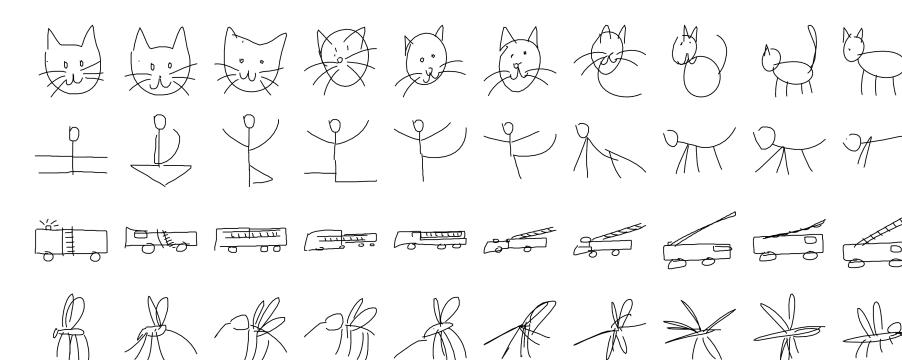
David Ha (hadavid@google.com), Douglas Eck (deck@google.com)



Motivation



"I never made a painting as a work of art. It's all research." – Pablo Picasso



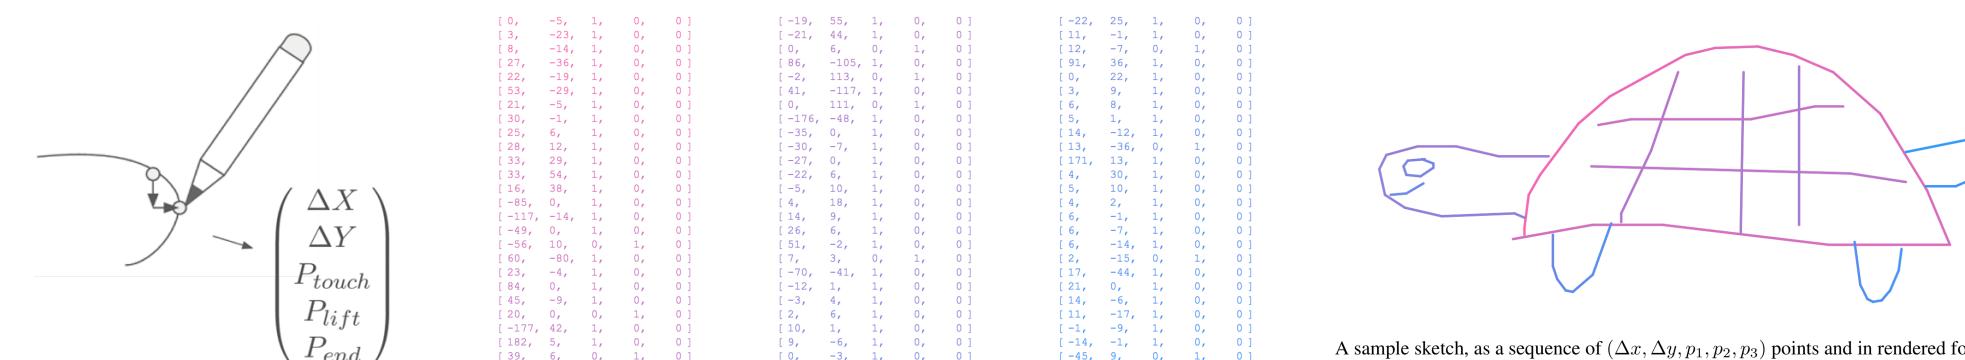
Humans, do not understand the world as a grid of pixels, but develop abstract concepts to represent what we see.

We learn to express a sequential, vector representation of an image as a short sequence of strokes.

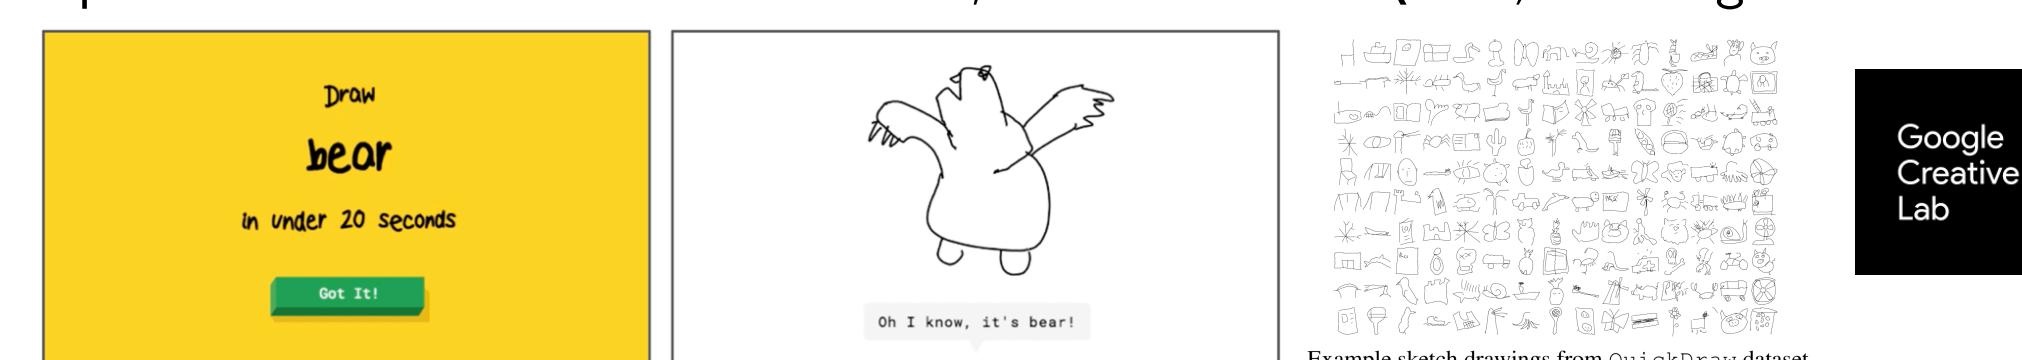
In this work, we investigate an alternative to pixel image modelling approaches, and propose a generative model for vector images.

Quick, Draw! Dataset

quickdraw.withgoogle.com/data

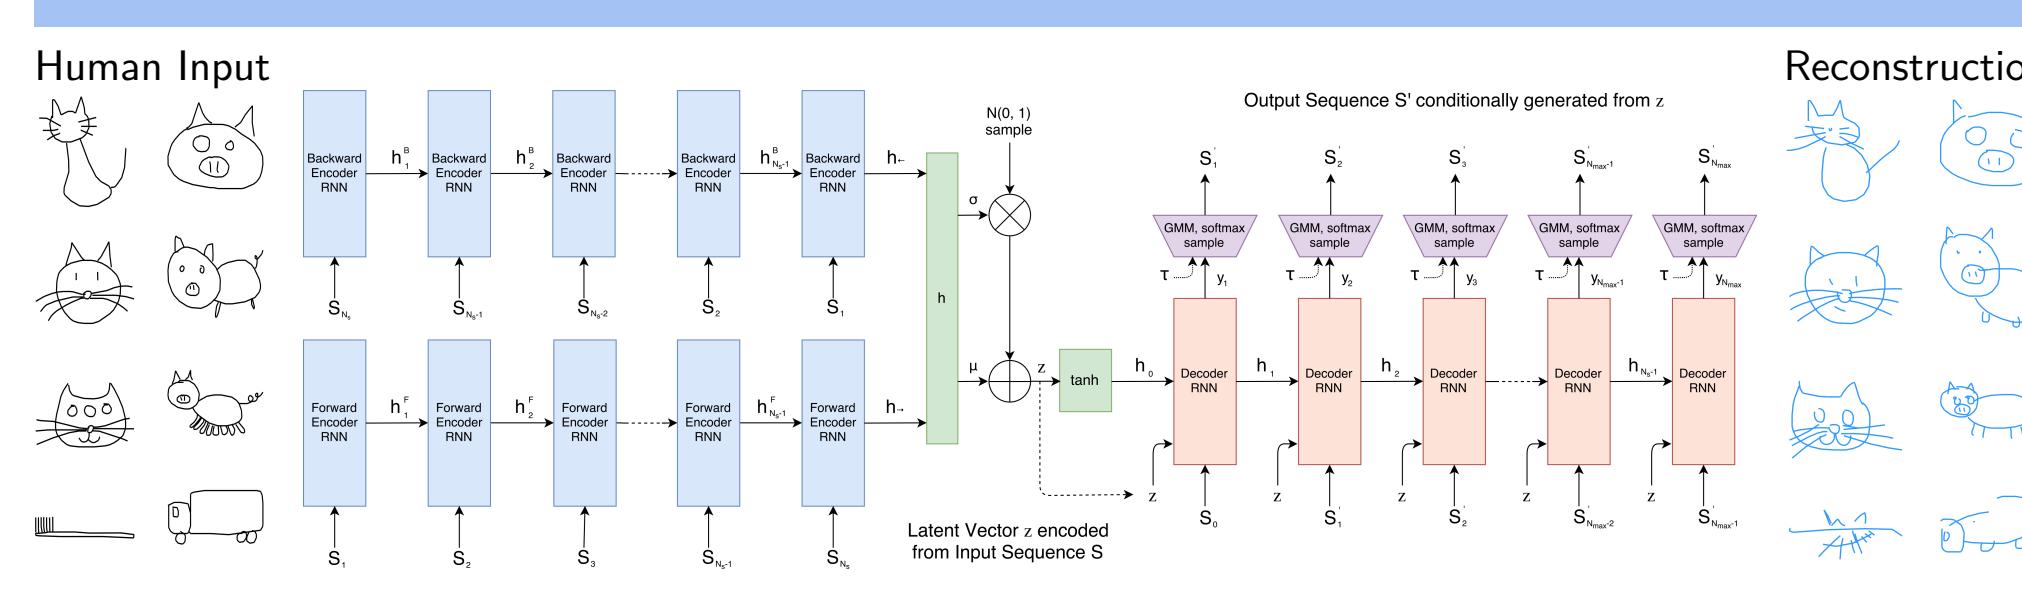


Sketches are represented as a sequence of motor actions controlling a pen. Open sourced dataset of 50M doodles, collected from Quick, Draw! game.

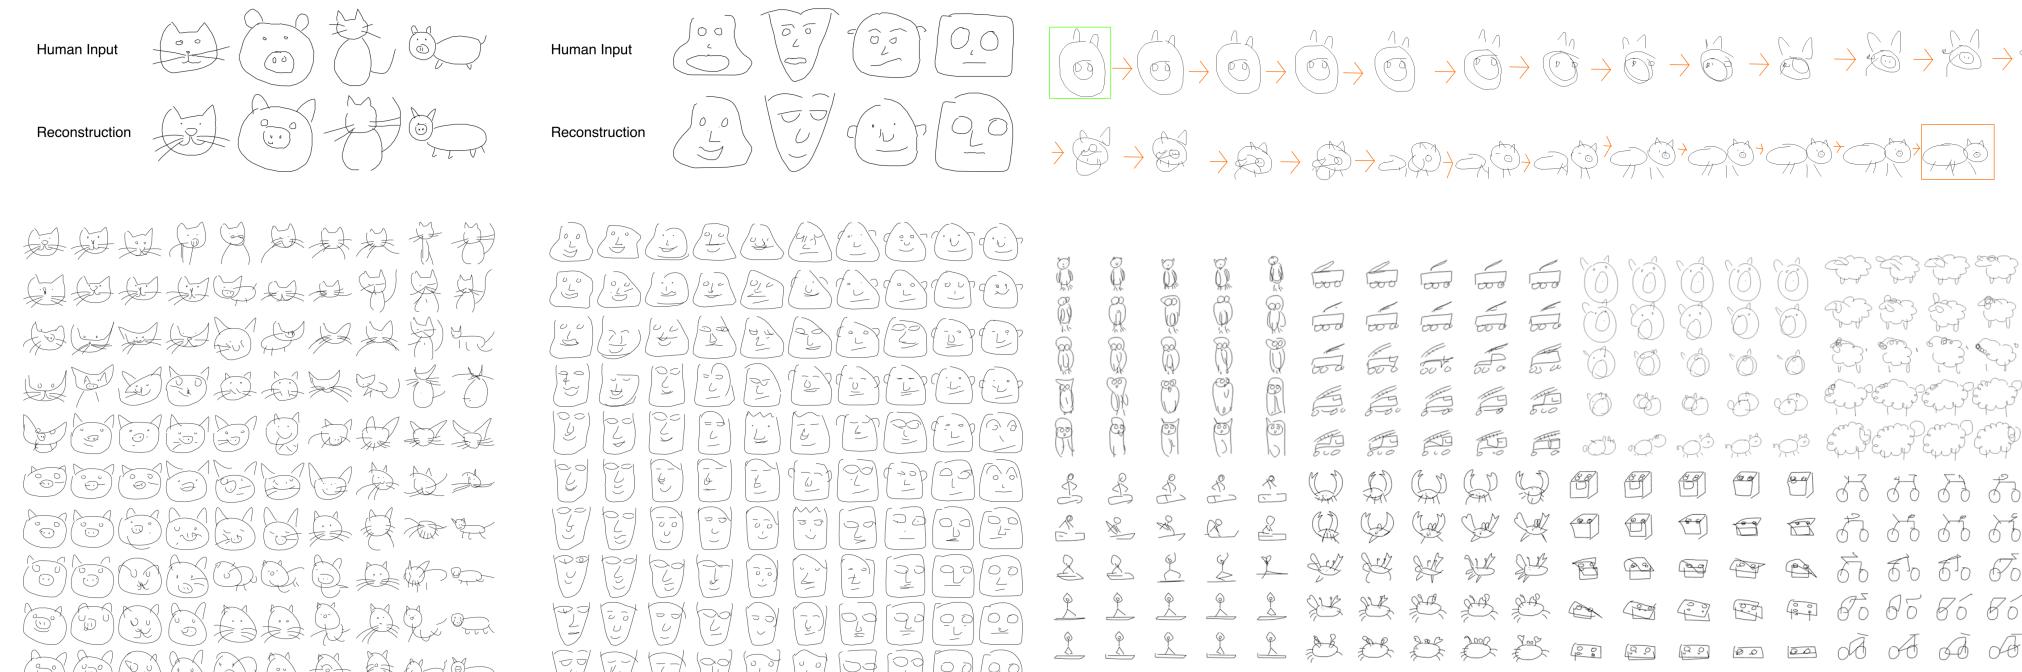


sketch-rnn

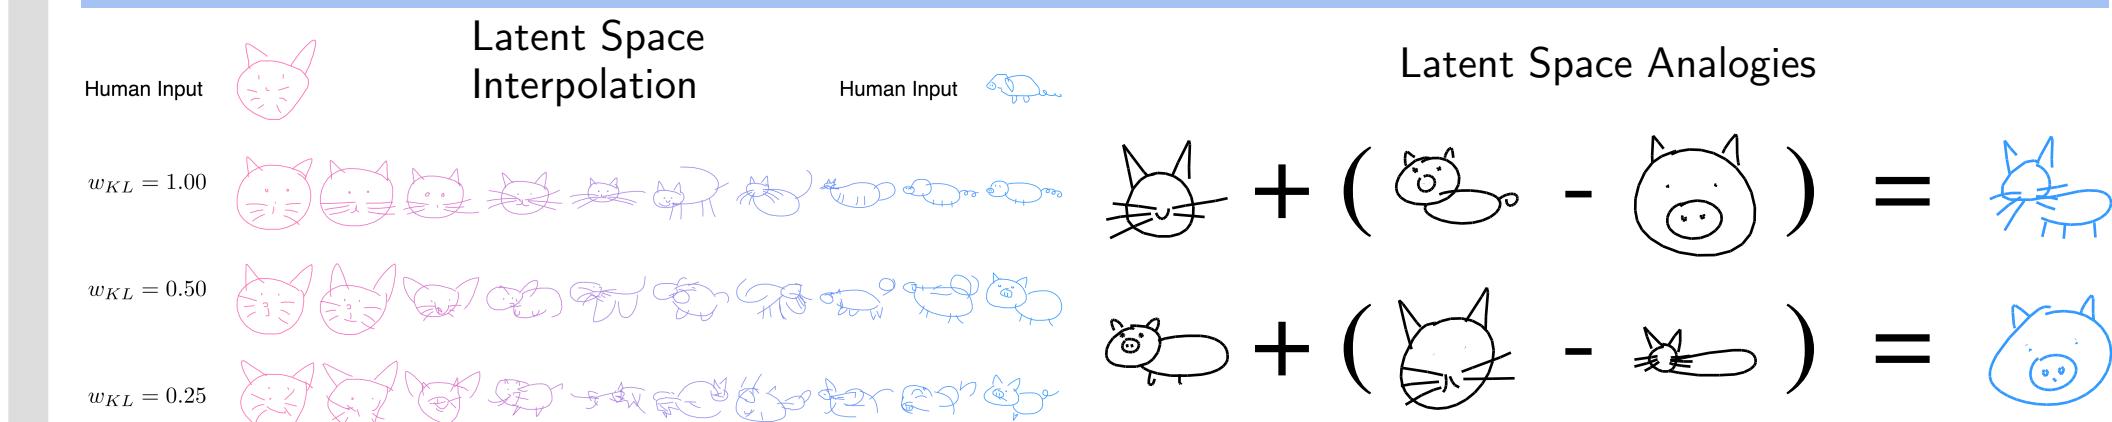
magenta.tensorflow.org/sketch-rnn-demo



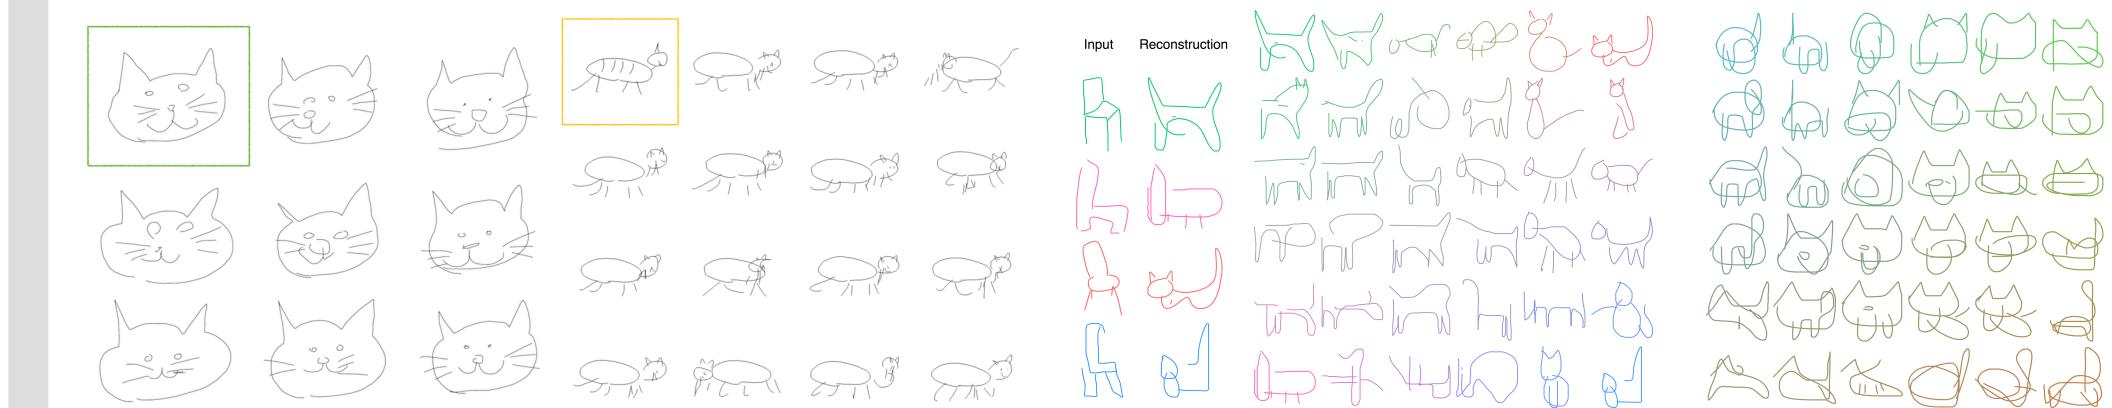
- Seq-to-Seq Variational Autoencoder
- Bidirectional RNN Encoder
- Mixture Density RNN Decoder
- Weighted Loss Function



Conditional Generation: Latent Space Exploration



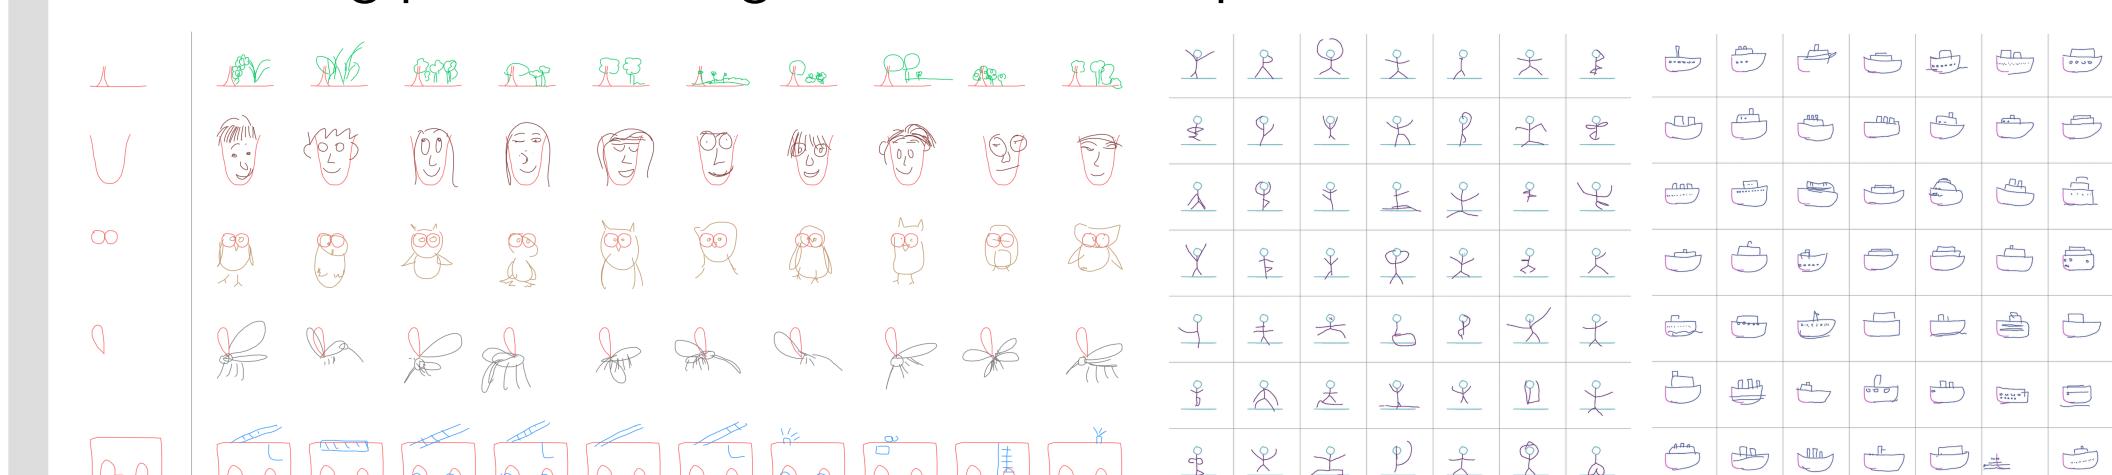
- Sample multiple similar drawings from a single human input sketch.
- Generate a drawing in the style of another class, and also interpolate for more.



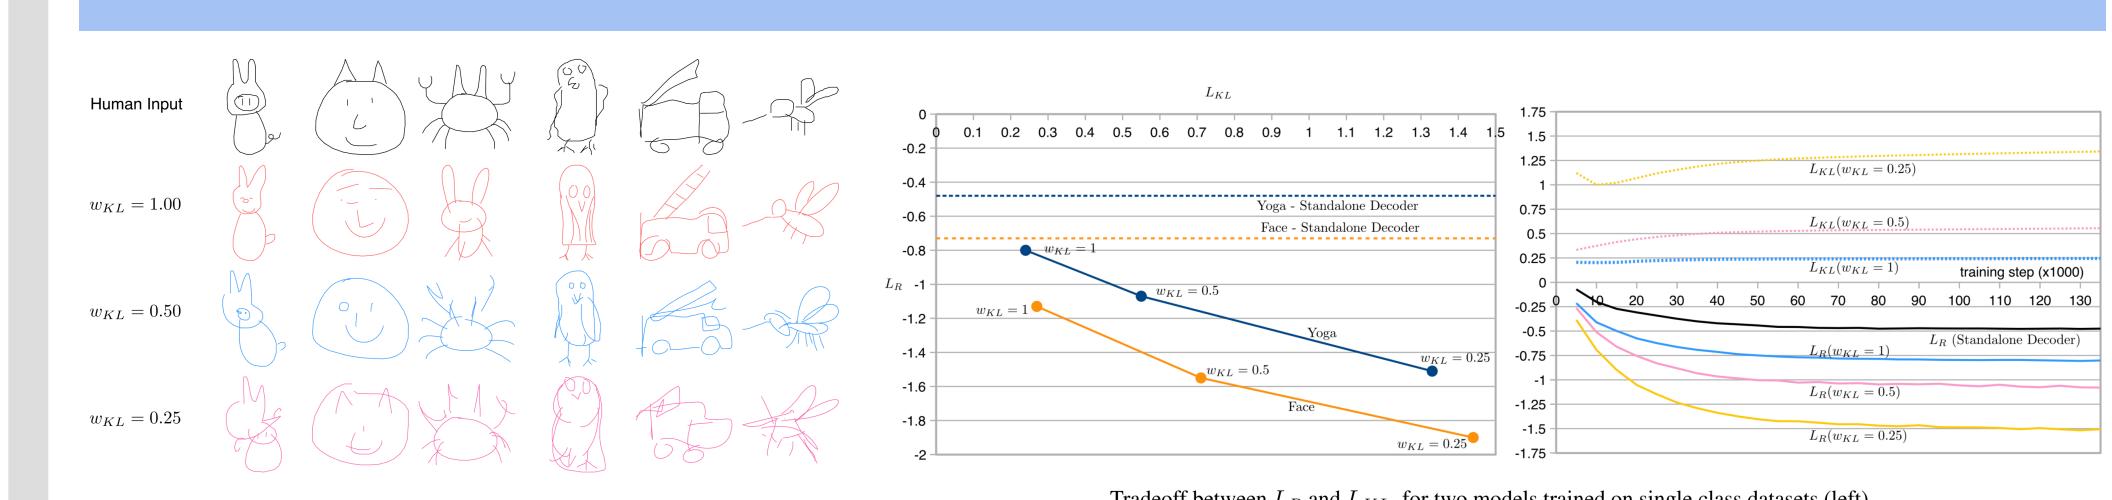
Unconditional Generation: Use Decoder Only



- Varying temperature parameter to control variability of generated sketches.
- Predicting possible endings of various incomplete sketches.



Which Loss Controls Image Coherency?



Other Datasets: Kanji, Aaron's 10K Sheep Market



Pixel → Strokes

- We can replace the encoder from a RNN to a ConvNet to read in pixels.
- Chen et al. (2017) extended sketch-rnn to use ConvNet encoder for pix2strokes:

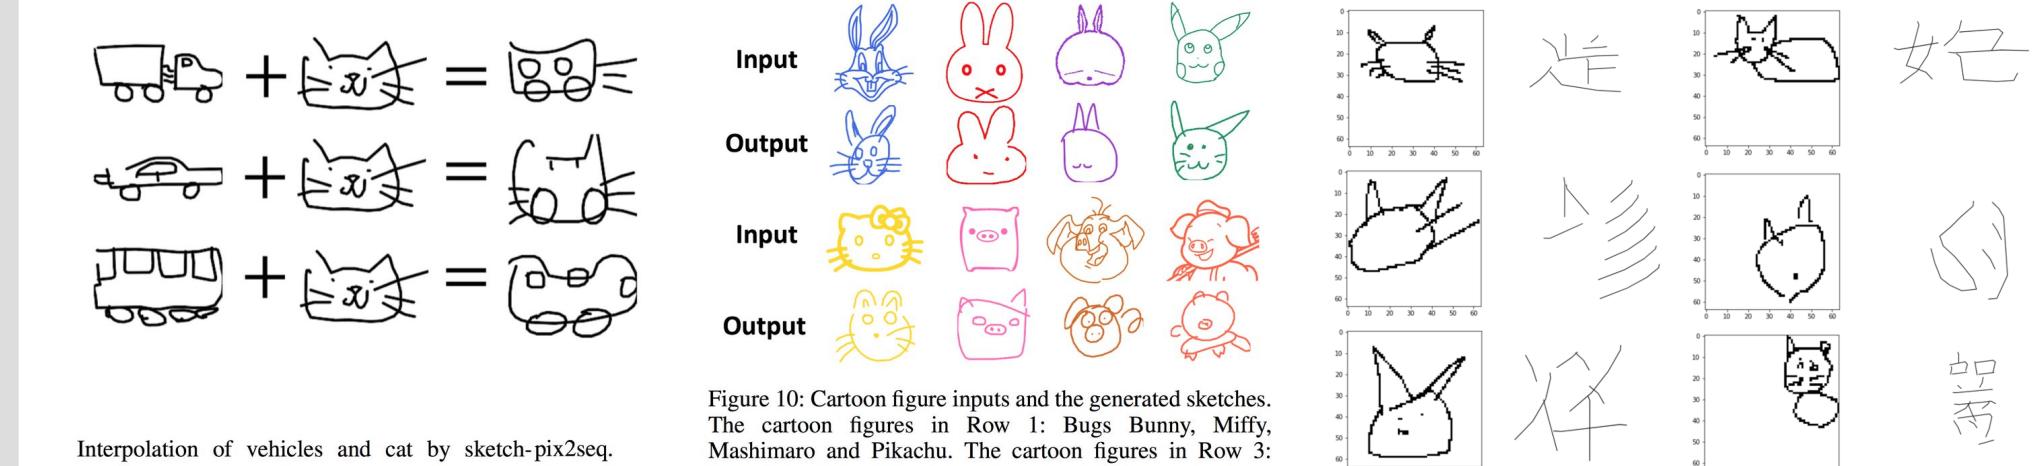
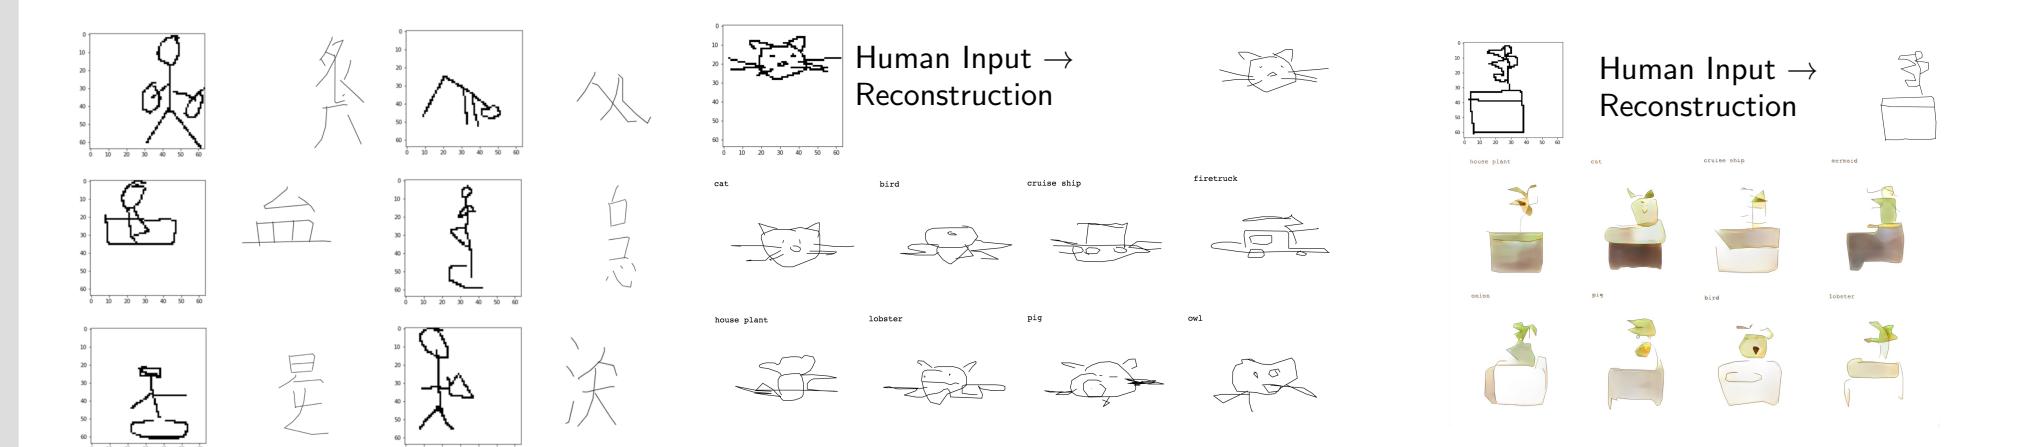


Figure 10: Cartoon figure inputs and the generated sketches. The cartoon figures in Row 1: Bugs Bunny, Miffy, Mashimashi and Pikachu. The cartoon figures in Row 3: Hello Kitty, Monuko Boo, Zhu Bajie and Zhu Bajie.

- We can try this technique across domains, such as QuickDraw → Kanji.



- Or convert a pixel image of one QuickDraw class into strokes of another class.

sketch-rnn with colorization

- Generated sketch drawings can be colorized using image-to-image translation.



References

- Alex Graves, Generating sequences with RNNs, 2013.
 Bowman et al., Generating Sentences from a Continuous Space, 2015.
 Zhang et al., Drawing and Recognizing Chinese Characters with RNN, 2016.
 Chen et al., Sketch-pix2seq: A Model to Generate Sketches of Multiple Categories, 2017.

Online Resources

- | | |
|------------------------|--|
| Fake Kanji Experiment | otoro.net/kanji |
| Quick, Draw! Game | quickdraw.withgoogle.com |
| Paints Chainer | github.com/pfnet/PaintsChainer |
| Paints Transfer | github.com/illyasviel/style2paints |
| sketch-rnn github repo | magenta.tensorflow.org/sketch_rnn |
| other datasets | github.com/hardmaru/sketch-rnn-datasets |

